

# Continual Re-Solving for HUNL Poker: Theoretical Backbone and Implementation Manual

## Abstract

This document consolidates the theoretical backbone of a practical, continual re-solving system for heads-up no-limit Texas Hold'em (HUNL) under a sparse action menu and depth-limited search. The goal is twofold: (i) to serve as an instruction manual for implementing the method end-to-end, and (ii) to present the elementary proofs used to verify core invariants in the accompanying work *HUNL\_AI* (the implementation paper). All other foundational results are cited to the established literature on continual re-solving (DeepStack). Formal statements and proofs written out below are restricted to those provided in *HUNL\_AI*; no new theory is introduced here. :contentReference[oaicite:0]index=0

## 1 Problem Setting and High-Level Design

**Game model.** HUNL is a two-player zero-sum extensive-form game with imperfect information, chance nodes (dealing), and alternating moves. Play proceeds across streets (preflop, flop, turn, river), with public-card revelation between betting rounds.

**Continual re-solving.** At each decision, re-solve the current public subtree to obtain a local strategy, act from that strategy, and carry forward only (a) the acting player's range and (b) the opponent's counterfactual value (CFV) vector, updated after own/chance/opponent transitions. Depth is limited for tractability; beyond the limit, leaf utilities are provided by learned CFV functions. (Foundational soundness and re-solving protocol are due to DeepStack; see fixed reference at the end.)

**Sparse actions and depth limits.** A first layer of {Fold, Call/Check, Pot, All-in} with lean replies balances strength and tractability; depth limits at the end of street on preflop and flop invoke CFV networks; the turn proceeds to terminal with exact endgame. These choices match the implementation/invariant suite described in *HUNL\_AI*. :contentReference[oaicite:1]index=1

## 2 Notation and Objects

- $K$  buckets (clusters) per street;  $r_1, r_2 \in \Delta^K$  denote bucketed ranges.
- $P > 0$  current pot;  $P_{\text{norm}} := P / (S_1^{(0)} + S_2^{(0)}) \in (0, 1]$ .
- $\varphi_{\text{board}} \in \{0, 1\}^{52}$  one-hot encoding of public cards.
- CFV networks (flop/turn) output per-bucket pot-fraction CFVs for both players.

### 3 CFV Networks and Outer Zero-Sum Layer

#### 3.1 Input/Output contract (instruction)

For each depth-limit query (end of peflop or flop), build

$$x = [\text{Pnorm} \mid \varphi_{\text{board}} \mid r_1 \mid r_2] \in \mathbb{R}^{1+52+2K},$$

and obtain unadjusted predictions  $(v_1, v_2) \in \mathbb{R}^K \times \mathbb{R}^K$  in *fractions of pot*. Multiply by  $P$  only if chip units are needed downstream. :contentReference[oaicite:2]index=2

#### 3.2 Per-sample outer zero-sum adjustment (statement and proof)

**Proposition 1** (Outer zero-sum identity). *Let  $s_1 = \langle r_1, v_1 \rangle$ ,  $s_2 = \langle r_2, v_2 \rangle$ , and  $\delta := \frac{1}{2}(s_1 + s_2)$ . Define*

$$f_1 := v_1 - \delta \mathbf{1}, \quad f_2 := v_2 - \delta \mathbf{1}.$$

*Then  $\langle r_1, f_1 \rangle + \langle r_2, f_2 \rangle = 0$  (up to numerical roundoff).*

*Proof.*

$$\langle r_1, f_1 \rangle + \langle r_2, f_2 \rangle = \langle r_1, v_1 \rangle - \delta \langle r_1, \mathbf{1} \rangle + \langle r_2, v_2 \rangle - \delta \langle r_2, \mathbf{1} \rangle = (s_1 + s_2) - \delta(1 + 1).$$

Since  $\delta = \frac{1}{2}(s_1 + s_2)$ , the expression reduces to  $(s_1 + s_2) - (s_1 + s_2) = 0$ . :contentReference[oaicite:3]index=3

□

**Implementation note.** Apply the adjustment samplewise after the linear heads and before loss computation; report all validation metrics in pot units.

### 4 Follow/Terminate Gadget and Monotone Carry-Forward

Let  $S$  be a boundary state (end-of-street) with opponent range  $r_{\text{opp}} \in \Delta^K$  and constraint vector  $c^{(t)} \in \mathbb{R}^K$  (opponent CFFs) carried from the previous boundary.

#### 4.1 Gadget (instruction)

For each bucket  $i$ , the opponent chooses *Terminate* (receive  $c_i^{(t)}$ ) or *Follow* (enter the re-solved subgame and receive  $\tilde{v}_i^{(t)}$ ). Define

$$v_{\text{opp},i} := \max\{c_i^{(t)}, \tilde{v}_i^{(t)}\}, \quad v_{\text{hero},i} := -v_{\text{opp},i}.$$

This preserves zero-sum per bucket and range-weights to zero overall (consistency with the outer layer). :contentReference[oaicite:4]index=4

#### 4.2 Monotone carry-forward (statement and proof)

**Proposition 2** (Protection & monotonicity). *With the update  $c^{(t+1)} := \max\{c^{(t)}, \tilde{v}^{(t)}\}$  (componentwise), it holds that*

(i)  $v_{\text{opp},i} \geq c_i^{(t)}$  for all  $i$  (Terminate is never worse than Follow).

(ii)  $c_i^{(t+1)} \geq c_i^{(t)}$  for all  $i$  (non-decreasing lower bound across boundaries).

*Proof.* (i) Since  $v_{\text{opp},i} = \max\{c_i^{(t)}, \tilde{v}_i^{(t)}\}$ ,  $v_{\text{opp},i} \geq c_i^{(t)}$  by definition. (ii) By the update,  $c_i^{(t+1)} = \max\{c_i^{(t)}, \tilde{v}_i^{(t)}\} \geq c_i^{(t)}$ . :contentReference[oaicite:5]index=5

□

**Implementation note.** Use self-play CFVs for  $c^{(t)}$  in practice for stability; the termination option guarantees safety while damping oscillations between boundaries. :contentReference[oaicite:6]index=6

## 5 Bayes-Consistent Range Updates

Let  $i \in \{1, \dots, K\}$  index opponent buckets. After observing an opponent action  $a$ , update

$$p'(i) = \frac{p(i) \pi(a | i)}{\sum_{j=1}^K p(j) \pi(a | j)}.$$

**Proposition 3** (Normalization and invariances). *For any prior  $p \in \Delta^K$  and likelihoods  $\pi(a | i) \geq 0$  with  $\sum_i p(i) \pi(a | i) > 0$ ,*

- (i)  $\sum_i p'(i) = 1$  (posterior is a probability vector);
- (ii) If  $\pi(a | i)$  is constant in  $i$ , then  $p' = p$  (uninformative action leaves the prior unchanged);
- (iii) For a sequence of independent observations  $a_1, \dots, a_T$ , the posterior equals the one-shot update with product likelihoods.

*Proof.* (i) Summing the numerator over  $i$  yields the denominator; division gives 1. (ii) If  $\pi(a | i) = \kappa$  for all  $i$ , then  $p'(i) = p(i)\kappa/(\kappa \sum_j p(j)) = p(i)$ . (iii) Follows by multiplication of likelihoods and a single normalization at the end. :contentReference[oaicite:7]index=7  $\square$

**Instruction.** Maintain strictly positive mass by construction; a small uniform mixture is acceptable as a safety fallback when external priors are degenerate. :contentReference[oaicite:8]index=8

## 6 Clustering and Bucket Interface

**Feature sketch.** Practical clustering can be framed in a low-dimensional feature space capturing (i) showdown equity, (ii) equity potential, and (iii) a payoff proxy, normalized before Euclidean  $k$ -means. Bucket count is  $K = \min\{K_{\text{target}}, |H|\}$  for candidate hand set  $H$ . Drift-triggered re-clustering avoids unnecessary recomputation. (Full engineering details are in the implementation paper; clustering is a design layer, not a theoretical guarantee, so no proof is replicated here.) :contentReference[oaicite:9]index=9

## 7 Depth-Limit Policy and Endgame

**Depth-limit placement.** End-of-street on preflop/flop (CFV networks) and solve-to-terminal from turn through river (exact endgame). This matches the documented interface between search and value functions and preserves the pot-fraction semantics. :contentReference[oaicite:10]index=10

## 8 Engine Invariants (Sanity Layer)

**Mass and pot monotonicity.** Range mass is conserved to numerical tolerance, and public pot changes are non-negative modulo explicit refunds, with street transitions restricted to “stay or advance by one”. (These are enforced as testable invariants rather than game-theoretic lemmas; no further proofs needed here.) :contentReference[oaicite:11]index=11

## 9 Implementation Checklist (Instruction Manual)

1. **State & legality.** Implement a single public-state engine: actor order per street; sparse legal menu; street advance on equalization; all-in lock fast-forward. :contentReference[oaicite:12]index=12
2. **CFV I/O.** Input  $[P_{\text{norm}} | \varphi_{\text{board}} | r_1 | r_2]$ ; output per-bucket pot-fraction CFVs; apply outer zero-sum adjustment before loss and logging. :contentReference[oaicite:13]index=13
3. **Re-solver.** CFR-style iterations with *Follow/Terminate* gadget at boundary; self-play constraints; carry-forward via componentwise max; statistical action pruning with reactivation (variance-aware). :contentReference[oaicite:14]index=14
4. **Preflop cache.** Content-addressed signature; require bit-identical reuse on hits. :contentReference[oaicite:15]index=15
5. **Sanity harness.** Enforce zero-sum residual  $\leq 10^{-6}$  (post-adjust), range-mass tolerance, and non-negative pot deltas on synthetic transitions. :contentReference[oaicite:16]index=16

## 10 References (Fixed)

- **Implementation paper:** *HUNL\_AI* (ResolveNet Poker: Neural CFV re-solving for HUNL; full engineering and verification suite). :contentReference[oaicite:17]index=17
- **Foundational continual re-solving:** Moravčík et al. (DeepStack: Expert-Level Artificial Intelligence in Heads-Up No-Limit Poker), *Science* 356(6337):508–513, 2017; see also arXiv:1701.01724 for the technical appendix (algorithms, action tables, zero-sum outer layer, depth-limit placement, and LBR settings).

## Appendix: Minimal Proof Sketches Included

Only proofs reproduced in *HUNL\_AI* are written out here: (A) outer zero-sum identity for the per-sample adjustment; (B) gadget protection and monotone carry-forward; (C) Bayes posterior normalization/invariance. All other theoretical foundations are cited to the fixed DeepStack reference above.